

SchwarzShuttle:

The Future of Taxi Operations

GCP Architecture for Real-Time Intelligence at Global Scale



Business Context & Challenges

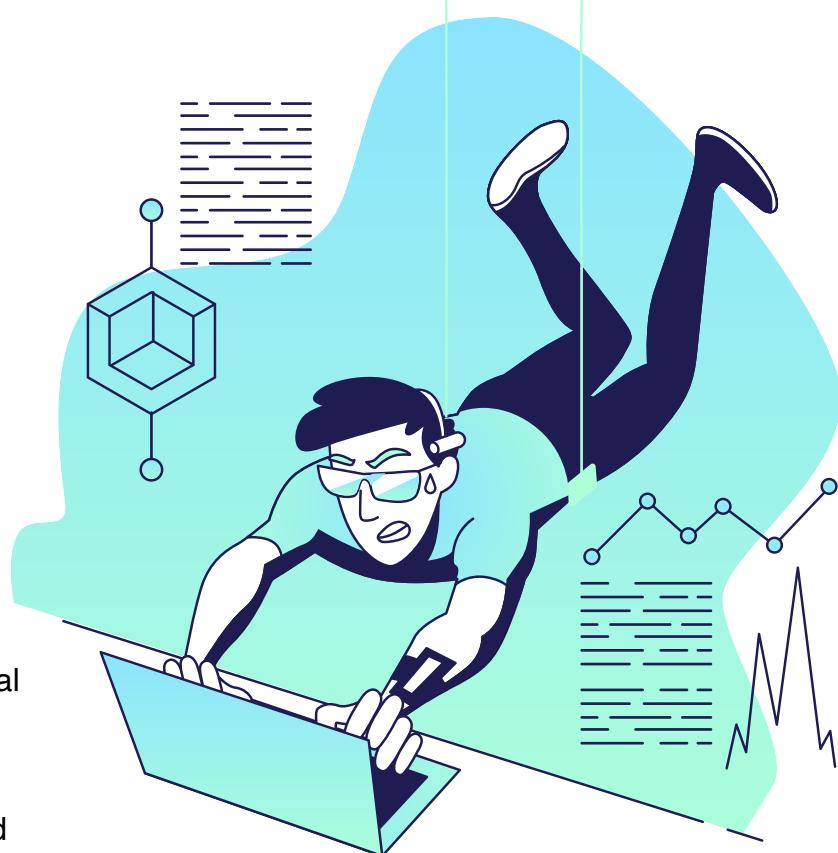
Our global taxi business operates at a massive scale, ensuring seamless transportation services worldwide.

❑ Operational Scale:

- Fleet of **11,000+ taxis** operating **24/7**
- Serving **millions of passengers** across multiple cities
- Integrating technology for real-time booking and tracking

❑ Key Challenges:

- **Data Management:** Handling vast amounts of trip, customer, and payment data efficiently
- **Security:** Ensuring the protection of sensitive user and financial data
- **Fraud Prevention:** Detecting and mitigating fraudulent transactions and unauthorised access
- **Optimisation:** Enhancing route planning, driver allocation, and fuel efficiency with AI-driven solutions



Solution Requirements

❑ Functional Requirements:

- Trip data collection and processing
- Real-time fleet monitoring and dispatch
- Revenue analysis and receipt generation
- Secure payment processing
- Customer communications management

❑ Non-Functional Requirements:

- **Performance:** Low latency for real-time operations
- **Scalability:** Support for global operations across time zones
- **Security:** Protection of sensitive customer and payment data
- **Reliability:** 24/7 availability with fault tolerance
- **Maintainability:** Component-level updates without service disruption



Data Ingestion & Processing Patterns



Trip Data

Start/end locations, distance, duration.

Processed via Pub/Sub → Cloud Functions → BigQuery

~500MB daily data volume across fleet

Payment Info

Fares, transaction logs, payment methods.

Processed with CMK encryption for PCI-DSS compliance

Secured via service-specific IAM roles and data governance



GPS Data

Real-time vehicle tracking, traffic conditions.

Streamed via Cloud IoT Core → Pub/Sub → Real-time analytics

High-frequency data (5-10 sec intervals) requiring scalable ingestion

Driver Ratings & Feedback

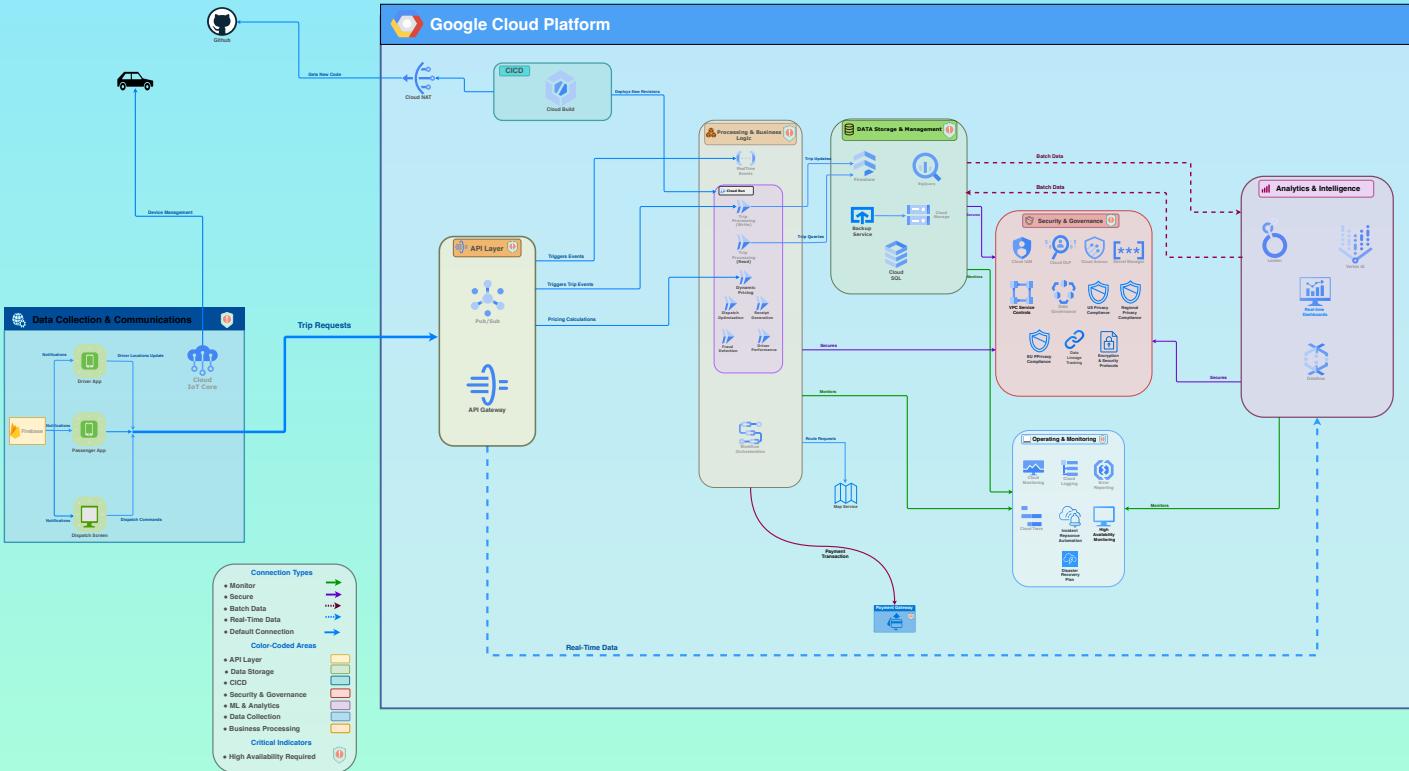
Service quality assessment.

Inputs to ML models for performance optimisation

Triggers operational alerts via Cloud Functions

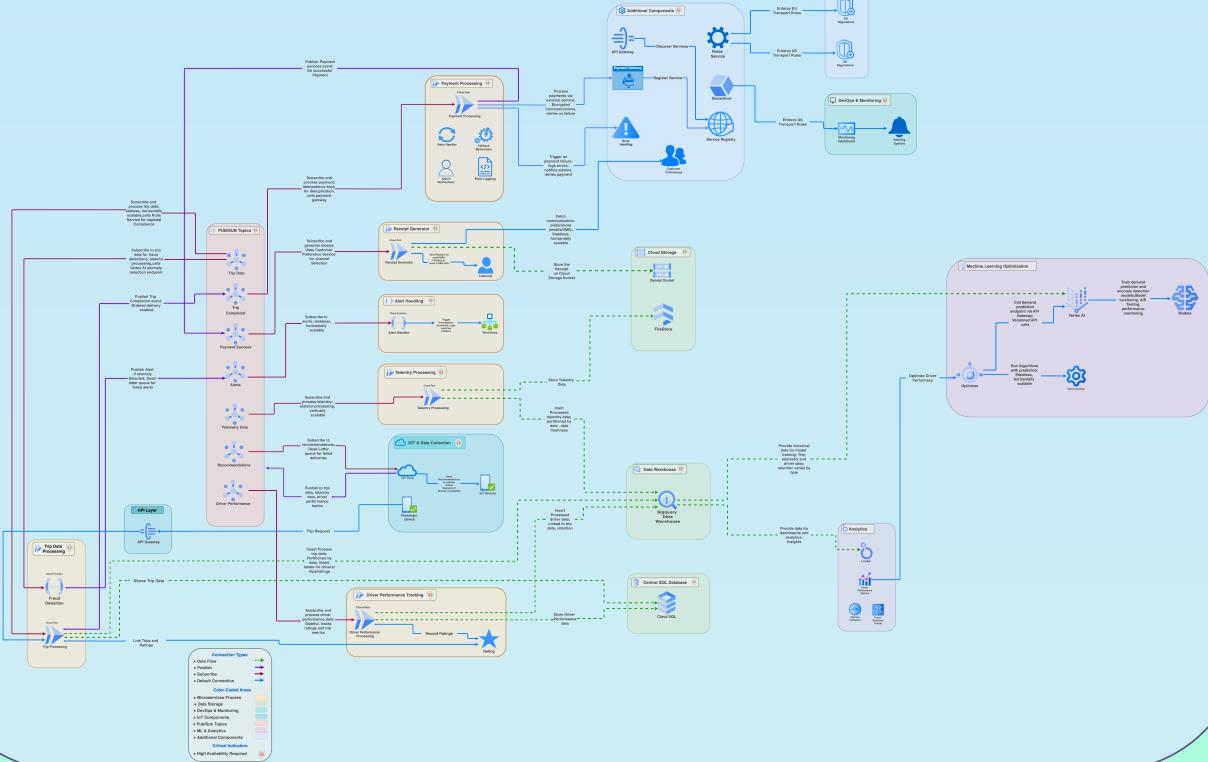


High-Level Architecture



View high-resolution diagram [here](#), and download drawio source file [here](#)

Detailed System Architecture



Key Points:

- Event-Driven Core:** Pub/Sub messaging enables real-time processing across all services
- Service Specialization:** Dedicated microservices for trip, payment, and telemetry workflows
- Data Pipeline Integration:** Seamless flow from operational systems to analytics
- ML Feedback Loop:** Continuous intelligence from data collection to prediction deployment
- Defense-in-Depth:** Security controls implemented at each architectural layer

View high-resolution diagram [here](#), and download drawio source file [here](#)

GCP Implementation



- | | | | | |
|--|--|---|--|--|
| Pub/Sub
(Real-Time Messaging & Event Streaming) <ul style="list-style-type: none">Enables real-time ingestion of IoT device data for fleet operations.Supports event-driven processing, triggering Cloud Functions for automation. | BigQuery
(Scalable Data Warehouse) <ul style="list-style-type: none">Stores and analyzes trip revenue, fleet performance, and fraud detection data.Uses CMEK encryption for secure data protection.Supports executive dashboards and analytics for decision-making. | Cloud Functions
(Serverless Compute for Automation) <ul style="list-style-type: none">Automates data processing, security monitoring, and fraud detection.Orchestrates workflows triggered by Pub/Sub events. | Vertex AI
(Machine Learning & Optimization) <ul style="list-style-type: none">Powers fleet optimization and fraud prevention models.Enables predictive maintenance for vehicles. | IAM & Security Tools
(Access Control & Threat Detection) <ul style="list-style-type: none">IAM enforces least privilege access to services.Log Sinks track security incidents in BigQuery.SonarQube, Trivy, OWASP ZAP integrate into CI/CD for security scans. |
|--|--|---|--|--|

Deployment Objectives



Security

- Protects sensitive trip and revenue data from cyber threats.
- Implements **Customer-Managed Encryption Keys (CMEK)** for secure storage.
- Enforces **IAM with least privilege access** to prevent unauthorised access.
- Integrates **log monitoring and security scans** (SonarQube, Trivy, OWASP ZAP) in CI/CD.



Scalability

- Supports **auto-scaling** to handle growing data from IoT devices.
- Ensures **low-latency** processing during peak fleet operations.
- Adapts to **multi-region** deployments for global expansion.



Optimization

- Enhances **real-time fleet management** using ML-powered insights.
- Reduces **operational costs** through intelligent resource allocation.
- Improves **user experience** with efficient and seamless data processing.

Automation Strategy



1

Dual Approach: Terraform (Primary) + Python (alternative)

2

Rationale:

- Terraform for declarative infrastructure
- Python for dynamic resource provisioning, complex logic

3

Key Capabilities:

- Component-level deployment
- Security configuration

4

Benefit: Hours → Minutes deployment time

Python Deployment Script

```
def create_bigquery_dataset(key_path_name):
    """Creates a BigQuery dataset with CMEK."""
    dataset = bigquery.Dataset(f"{PROJECT_ID}.{DATASET_ID}")
    dataset.default_encryption_configuration = bigquery.EncryptionConfiguration(kms_key_name=key_path_name)
    try:
        bq_client.create_dataset(dataset)
        logger.info(f"Created dataset: {DATASET_ID} with CMEK")
    except Exception as e:
        logger.warning(f"Dataset exists or error: {e}")

def deploy_cloud_functions():
    """Deploys placeholder Cloud Functions for trip processing, receipt generation, and alerts."""
    functions = [
        ("trip_processor", "trip_data-sub", "process_trips.py"),
        ("receipt_generator", "trip_completed-sub", "generate_receipts.py"),
        ("alert_handler", "alerts-sub", "handle_alerts.py")
    ]
    for name, trigger_sub, source_file in functions:
        function_path = functions_client.cloud_function_path(PROJECT_ID, LOCATION, name)
        try:
            functions_client.create_function(request={
                "location": f"projects/{PROJECT_ID}/locations/{LOCATION}",
                "function": {
                    "name": function_path,
                    "source_archive_url": f"gs://{PROJECT_ID}-functions/{source_file}",
                    "entry_point": name,
                    "event_trigger": {
                        "event_type": "google.pubsub.topic.publish",
                        "resource": pubsub_publisher.topic_path(PROJECT_ID, trigger_sub.split('-')[0])
                    },
                    "runtime": "python39"
                }
            })
            logger.info(f"Deployed Cloud Function: {name}")
        except Exception as e:
            logger.warning(f"Function {name} exists or error: {e}")
```

Key Points:

- **Component-Level Deployment:** Each service deployed independently
- **Event-Driven Architecture:** Functions triggered by Pub/Sub events
- **Idempotent Operations:** Includes error handling for resilience
- **Configuration Flexibility:** Customisable runtime and locations
- **Audit Logging:** Deployment actions tracked for accountability

Terraform Infrastructure

Infrastructure-Terraform-main

Name

assets

- buckets.tf
- cloud-run.tf
- cloudsql.tf
- datasets.tf
- firestore.tf
- kms-keys.tf
- main.tf
- network.tf
- pub-sub.tf
- README.md
- service-accounts.tf

vars.yaml

Key Points:

- Configuration as Code:** Complete infrastructure defined in version control
- Modular Architecture:** Separated services with dedicated resources
- Security-First Design:** Private networking and service-specific accounts
- Dynamic Scaling:** Automatic resource adjustment based on demand
- External Configuration:** YAML-driven deployment for environment flexibility

```
# Cloud Run services with secure configuration
resource "google_cloud_run_v2_service" "processing" {
  for_each      = local.yaml_vars.cloud-runs
  name          = each.value.name
  location      = local.yaml_vars.public-vars.region
  deletion_protection = each.value.deletion_protect
  ingress        = each.value.ingress

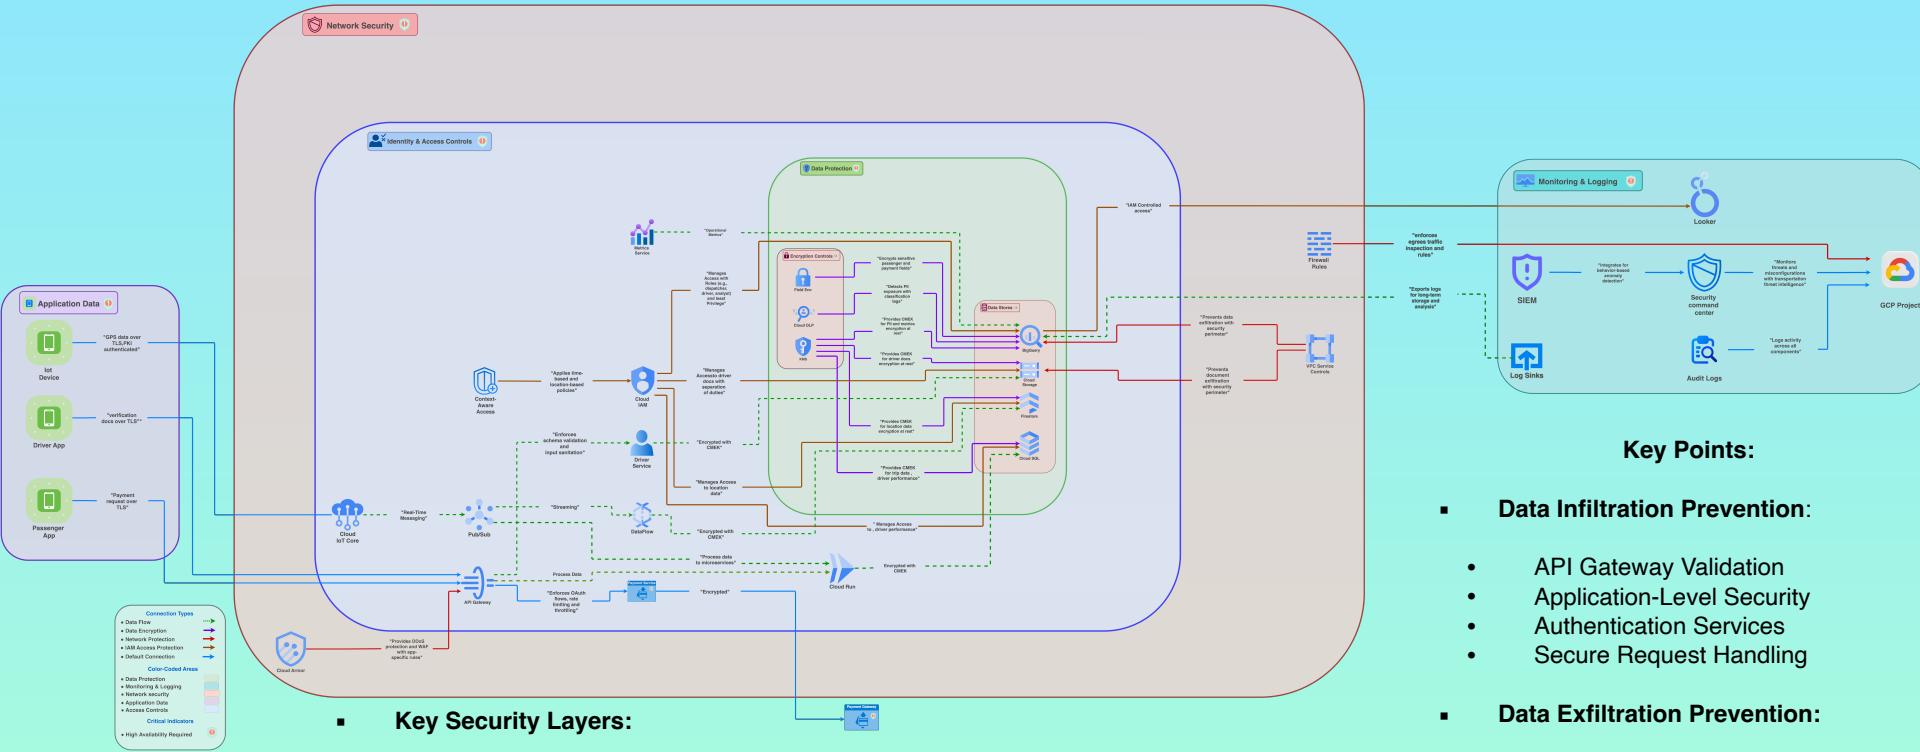
  template {
    service_account = google_service_account.service_accounts[each.value.service_account_key].email
    scaling {
      max_instance_count = each.value.max_count
    }
    containers {
      image = each.value.image
      resources {
        limits = {
          cpu   = each.value.cpu
          memory = each.value.memory
        }
      }
    }
  }
}

# Secure networking infrastructure
resource "google_compute_network" "vpc_network" {
  name          = local.yaml_vars.network.network_name
  auto_create_subnetworks = false
}

resource "google_compute_subnetwork" "private_subnet" {
  name          = "${local.yaml_vars.network.network_name}-private"
  ip_cidr_range = local.yaml_vars.network.private_subnet_cidr
  region        = local.yaml_vars.public-vars.region
  network       = google_compute_network.vpc_network.id
}
```

Browse complete repository [here](#)

Multi-Layered Security & Data Protection Strategy



▪ Key Security Layers:

- Perimeter Security: VPC, Firewalls
 - Identity & Access: IAM with Least Privilege
 - Data Protection: CMK Encryption
 - Application Security: Input Validation, Secure APIs

- **Comprehensive Protection:**

- Mitigating External Threats
 - Managing Insider Risks
 - End-to-End Security Approach

Key Points:

Data Infiltration Prevention:

- API Gateway Validation
 - Application-Level Security
 - Authentication Services
 - Secure Request Handling

Data Exfiltration Prevention:

- Network Perimeter Controls
 - Restricted Data Movement
 - Authorized Access Only
 - Comprehensive Logging & SIEM Monitoring

*View high-resolution
diagram [here](#), and download
drawio source file [here](#)*

Security Implementation Controls

Python: Data Encryption & Access Control:

```
def create_kms_key():
    """Creates a KMS keyring and key for encryption."""
    location = "global"
    keyring_path = kms_client.key_ring_path(PROJECT_ID, location, KEYRING_NAME)
    try:
        kms_client.create_key_ring(
            request={"parent": f"projects/{PROJECT_ID}/locations/{location}",
                     "key_ring_id": KEYRING_NAME})
        logger.info(f"Created keyring: {KEYRING_NAME}")
    except Exception as e:
        logger.warning(f"Keyring exists or error: {e}")

    key_path = kms_client.crypto_key_path(PROJECT_ID, location, KEYRING_NAME, KEY_NAME)
    try:
        kms_client.create_crypto_key(request={
            "parent": keyring_path,
            "crypto_key_id": KEY_NAME,
            "crypto_key": {"purpose": kms_v1.CryptoKey.CryptoKeyPurpose.ENCRYPT_DECRYPT}
        })
        logger.info(f"Created key: {KEY_NAME}")
    except Exception as e:
        logger.warning(f"Key exists or error: {e}")
    return key_path

def setup_security_monitoring():
    """Configures log sinks to BigQuery for monitoring."""
    # Configure audit logging
    client = gc.logging.Client(credentials=credentials)
    sink_name = "schwarzshuttle_logs"
    sink = client.sink(sink_name)

    # Filter for security-relevant logs
    sink.filter_ = f"logName:'projects/{PROJECT_ID}/logs/cloudaudit.googleapis.com%'"
    sink.destination = f"bigquery.googleapis.com/projects/{PROJECT_ID}/datasets/{DATASET_ID}/tables/audit_logs"

    sink.create()
    logger.info(f"Created Log Sink '{sink_name}' to BigQuery.")
```

Terraform: Network & Security Configuration

```
# Secure service identity management
resource "google_service_account" "service_accounts" {
    for_each  = local.yaml_vars.service_accounts
    account_id = each.value.name
    display_name = each.value.display_name
}

# Customer-managed encryption keys
resource "google_kms_key_ring" "key_ring" {
    name      = local.yaml_vars.kms-keys.keyring.name
    location  = local.yaml_vars.public-vars.region
}

resource "google_kms_crypto_key" "crypto_key" {
    for_each  = local.yaml_vars.kms-keys
    name      = each.value.name
    key_ring = google_kms_key_ring.key_ring.id
}

# Encrypted data storage with fine-grained access control
resource "google_bigquery_dataset" "dataset-test" {
    for_each  = local.yaml_vars.Datasets
    dataset_id = each.value.dataset_id
    location  = local.yaml_vars.public-vars.region
    depends_on = [google_kms_crypto_key.crypto_key]

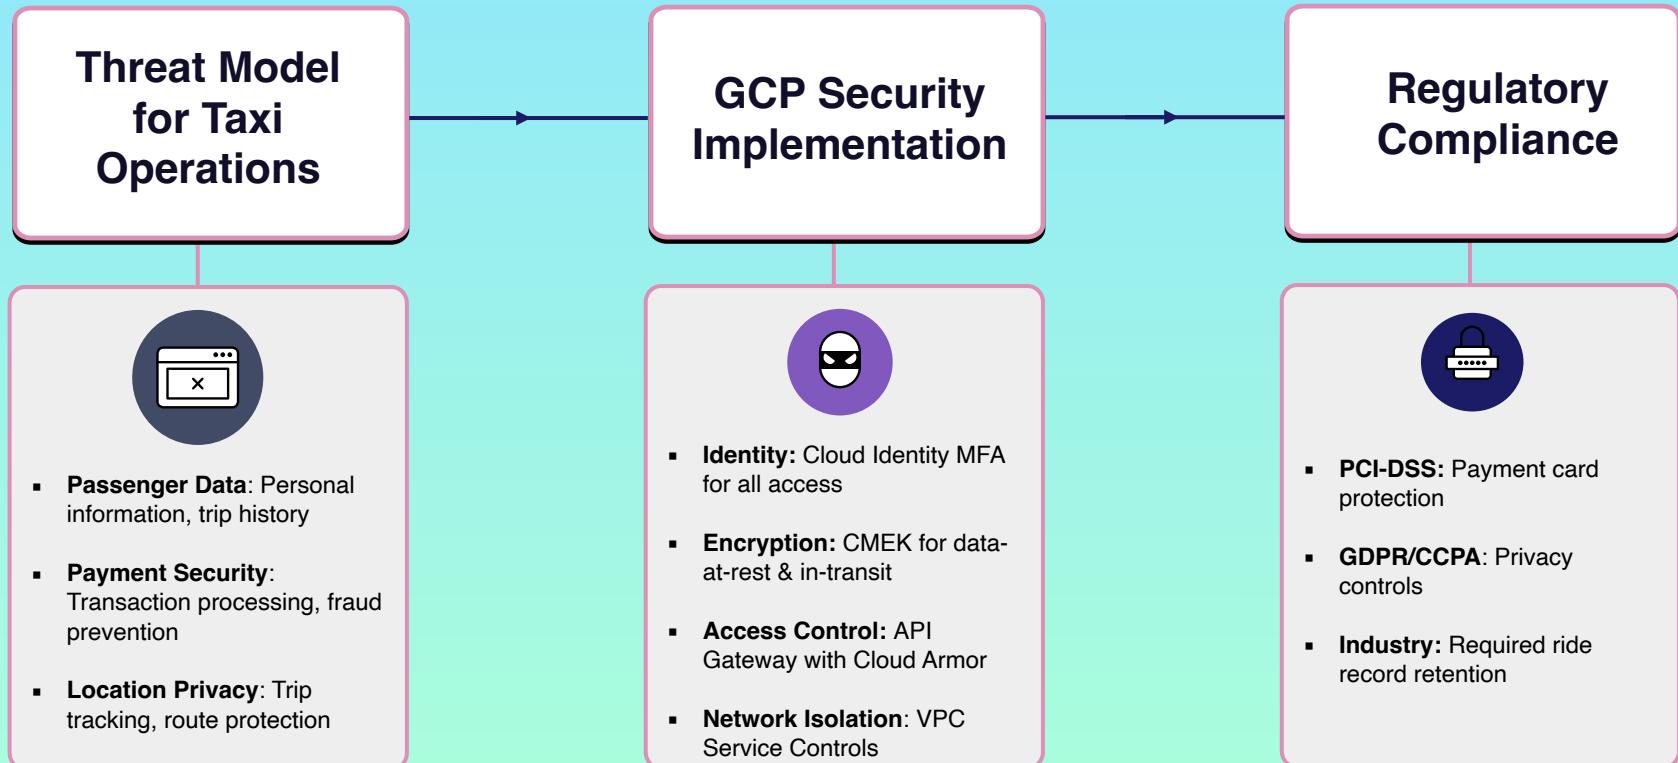
    default_encryption_configuration {
        kms_key_name = local.yaml_vars.kms-keys.dataset-key.name
    }
}

# Role-based access control for data
resource "google_bigquery_dataset_iam_binding" "data-reader" {
    for_each  = local.yaml_vars.Datasets
    dataset_id = google_bigquery_dataset.dataset-test[each.key].dataset_id
    role      = "roles/bigquery.dataViewer"
    members   = each.value.reader_members
}
```

Key Security mechanisms demonstrated:

- Customer-Managed Encryption Keys
- Least Privilege Access Control
- Fine-Grained IAM Bindings
- Security Monitoring
- Resource Isolation

Security Enhancement



GCP Defense-in-Depth for Taxi Operation

Monitoring & Incident Response

*Security Command Center
Cloud Logging
Event Threat Detection*

Endpoint & Identity Security

*Cloud Identity
IAM service accounts
Conditional access*

Data Security

*CMEK encryption
Cloud KMS
Cloud DLP*

Application Security

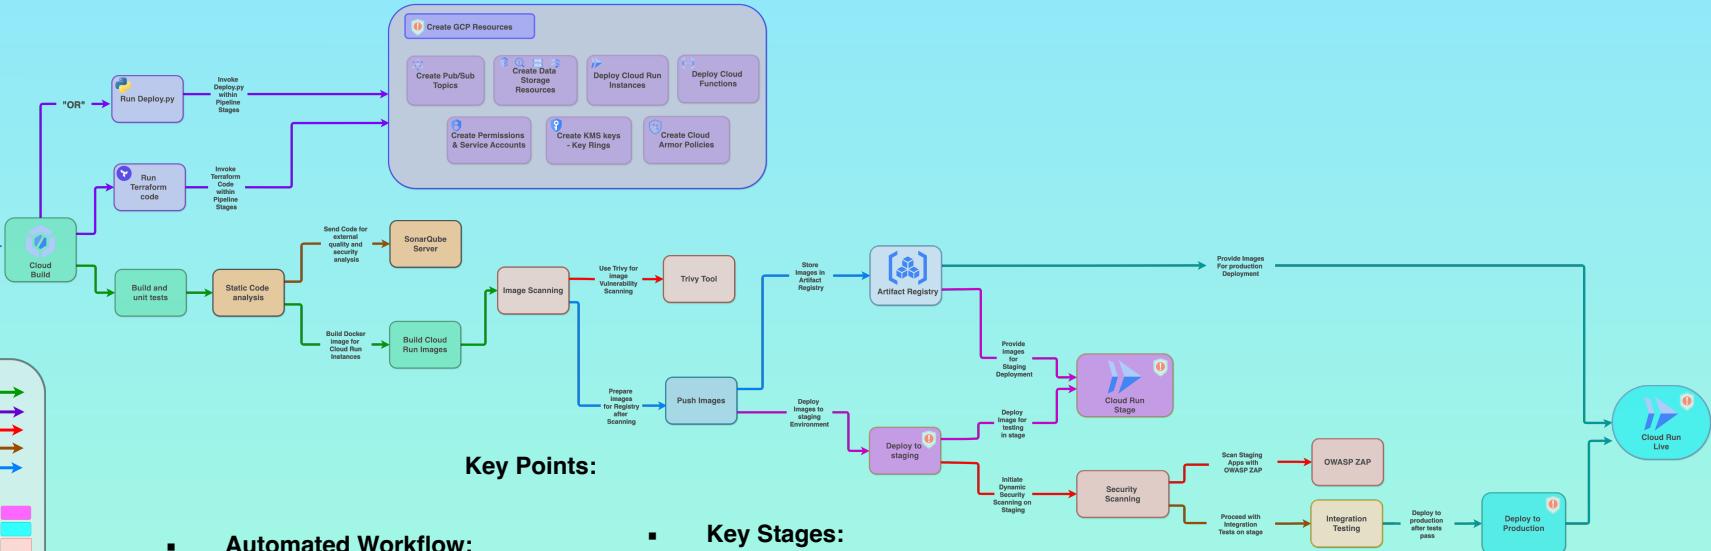
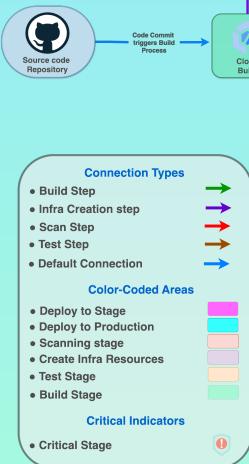
*Cloud Armor
API Gateway
Container security*

Perimeter Security (Network Level)

*VPC Service Controls
Cloud NAT
Secure IoT connectivity*



DevSecOps Pipeline



Key Points:

- Automated Workflow:**
 - Source Code to Cloud Deployment
 - Continuous Integration & Delivery
- Security Integration:**
 - Image Scanning
 - OWASP ZAP Security Checks
 - Vulnerability Assessment
- Key Stages:**
 - Build & Test
 - Security Scanning
 - Cloud Deployment
 - Artifact Management
- Advanced Features:**
 - Cloud Resource Provisioning
 - Automated Security Policies
 - Multi-Environment Support

View high-resolution diagram [here](#), and download drawio source file [here](#)

Measurable Improvements



Operational Metrics:

- 30% reduction in passenger wait times
- 15% decrease in vehicle idle time
- 8% increase in revenue per mile



Technical Metrics:

- 50% faster deployment time
- 60% reduction in configuration errors
- 25% reduction in security incidents
- 30% less service disruption during updates

Our Journey: From Launch to Leadership (2025-2027)

Year 1

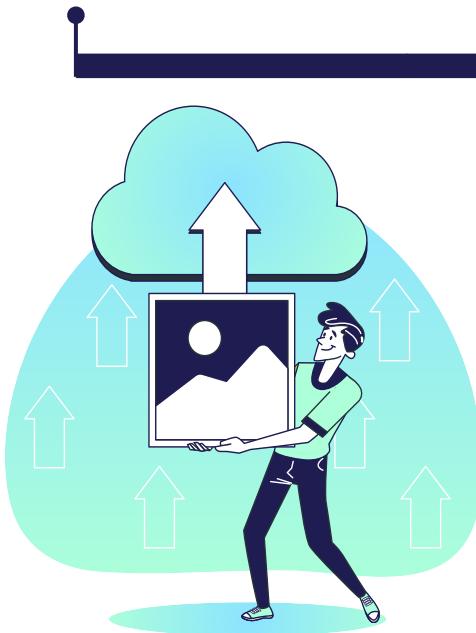
- Enhanced Analytics
- Infrastructure Improvements
- Security Advancements

Year 2

- Predictive Intelligence
- Customer Experience
- Operational Efficiency

Year 3

- Autonomous Capabilities
- Sustainability Focus
- Advanced Analytics

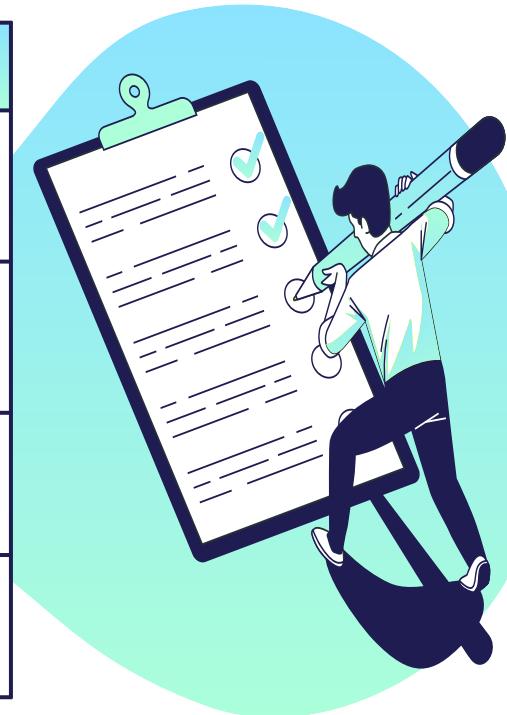


Expected Business Outcomes:

- **Higher utilization rates** (less idle time).
- **Better customer retention** (faster rides, optimized pricing).
- **Cost reduction** (fuel savings, route efficiency).

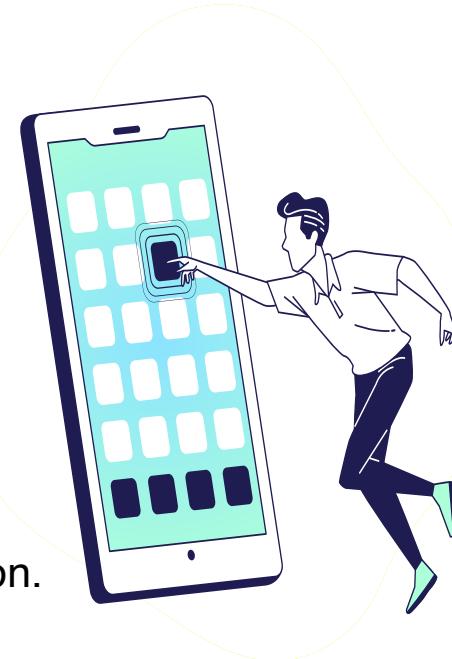
Requirements Fulfilled

Requirement	Implementation	Benefits
Architecture	Complete GCP solution with real-time processing and bidirectional communication	Comprehensive data management, operational control, actionable insights
Automation	Terraform (primary) + Python script (alternative) with component-level updates	Consistent deployments, operational flexibility, faster updates
Security	Defense-in-depth with CMEK, IAM, monitoring	Protected customer data, regulatory compliance, threat prevention
Vision	Strategic 3-year roadmap (2025-2027)	Clear evolution path aligned with business growth



Key Takeaways

- **Comprehensive Architecture** – End-to-end solution for global taxi operations.
- **Flexible Automation** – Terraform + Python options for deployment preferences.
- **Security By Design** – Multi-layered protection for sensitive data.
- **Measurable Impact** – Clear improvements in key business metrics.
- **Strategic Vision** – Forward-looking roadmap for continued innovation.



Driving the Future of Intelligent Mobility

From Data to Movement, From Insights to Action

